

# Multi-node Deployment of a Stateless Application on Minikube

---

This folder contains the configuration for the deployment and services of a web application that exposes everything via Ingress.

## Explanation

The deployment creates 2 pod replicas with a Hello-World container. The Service then groups all the pods and exposes them with a ClusterIP. Finally, the Ingress exposes the service externally, always on port 80.

The affinity section prevents the pods from being created on the same node at deployment time. Note that with `requiredDuringSchedulingIgnoredDuringExecution` it is allowed for the pods to end up on the same node at a later time, for example after a restart.

## Usage

### Prerequisites

- Minikube installed and working
- Kubectl installed and working
- Ingress addon enabled

### Starting Minikube

To start Minikube with multiple nodes, run the following command:

```
minikube start --nodes 2 --driver=docker -p multinode-demo
```

A Minikube profile is also created (not sure what it's for, but the tutorial included it). You can verify the presence of the nodes with the command:

```
kubectl get nodes
```

and

```
minikube status -p multinode-demo
```

### Deployment

To create the cluster, run the following command:

```
kubectl apply -f manifests/
```

## Verifying the Deployment

To check that the deployment was successful, run the following command:

```
kubectl get pods -o wide
```

and use other commands to check services and ingress.

## Accessing the Service

To access the service, it is not enough to connect to the cluster IP because it is not exposed.

```
curl http://192.168.49.2:31000
```

In fact, we need to open the tunnel for the Ingress.

```
minikube tunnel -p multinode-demo
```

After that, you can access the service with the command:

```
curl --resolve "www.example.com:80:127.0.0.1" -i http://www.example.com
```

## Cleaning up

To clean up use the `cleanup.sh` in the `scripts` directory:

```
./cleanup.sh
```